

Fraud Detection - Credit Card

About Dataset

The dataset has a time column which shows the transaction in seconds. The dataset have more columns from V1 to V28 which represents some feature about the transaction but as the transactions of credit cards are sensitive, the columns are presented by numbers. Moreover the data has 'Amount' column which is shows the transaction amount in dollars and 'Class' column shows whether the transaction is fraud or not.

Kaggle link for the Dataset - <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>
(<https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>).

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
credit_card = pd.read_csv('creditcard.csv')
```


In [3]:

```
credit_card.head()
```

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

5 rows × 31 columns



In [4]:

```
credit_card.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null float64
1   V1          284807 non-null float64
2   V2          284807 non-null float64
3   V3          284807 non-null float64
4   V4          284807 non-null float64
5   V5          284807 non-null float64
6   V6          284807 non-null float64
7   V7          284807 non-null float64
8   V8          284807 non-null float64
9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

In [5]:


```
credit_card.isnull().sum()
```

Out[5]:

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
```

In [6]:

```
sns.heatmap(credit_card.isnull())
plt.show()
```



In [7]:

```
cc = credit_card.copy()
```

In [8]:

```
cc.head()
```

Out[8]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

5 rows × 31 columns

In [9]:

```
cc['Class'].value_counts()
```

Out[9]:

```
0    284315
1      492
Name: Class, dtype: int64
```

In [10]:

```
cc.groupby('Class').mean()
```

Out[10]:

	Time	V1	V2	V3	V4	V5	V6	V7
Class								
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.002419	0.009637
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731

2 rows × 30 columns

In [11]:

```
normal = cc[cc.Class == 0]
fraud = cc[cc.Class == 1]
```

In [12]:

```
normal.shape , fraud.shape
```

Out[12]:

```
((284315, 31), (492, 31))
```

In [13]:

```
normal.Amount.describe()
```

Out[13]:

```
count      284315.000000
mean         88.291022
std        250.105092
min           0.000000
25%          5.650000
50%         22.000000
75%         77.050000
max       25691.160000
Name: Amount, dtype: float64
```

In [14]:

```
fraud.Amount.describe()
```

Out[14]:

```
count      492.000000
mean       122.211321
std       256.683288
min           0.000000
25%          1.000000
50%          9.250000
75%        105.890000
max       2125.870000
Name: Amount, dtype: float64
```

In [15]:

```
#As the data is not distributed evenly, we must try to disribute evenly
normal_sample = normal.sample(n = 492)
```

In [16]:

```
cc_final = pd.concat([normal_sample, fraud], axis = 0)
```

In [17]:

```
cc_final.sample(5)
```

Out[17]:

	Time	V1	V2	V3	V4	V5	V6	V7
238127	149534.0	1.984559	-1.930051	-1.088922	-1.675583	-1.137633	0.197087	-1.172027
86001	61038.0	-1.120009	0.750977	2.561013	-0.030162	-0.294961	0.376599	0.342202
18773	29753.0	0.269614	3.549755	-5.810353	5.809370	1.538808	-2.269219	-0.824203
135718	81372.0	-0.885254	1.790649	-0.945149	3.853433	-1.543510	0.188582	-2.988383
53794	46149.0	-1.346509	2.132431	-1.854355	2.116998	-1.070378	-1.092671	-2.230986

5 rows × 31 columns

In [18]:


```
#Checking if the data is evenly distributed or not
dis = cc_final['Class'].value_counts()
dis
```

Out[18]:

```
0    492
1    492
Name: Class, dtype: int64
```

In [19]:

```
dis.plot(kind = 'bar')
plt.show()
```



In [20]:

```
cc_final.groupby('Class').mean()
```

Out[20]:

	Time	V1	V2	V3	V4	V5	V6	V7
Class								
0	94732.664634	-0.085919	-0.019114	-0.053619	0.006691	-0.050785	-0.012492	0.024236
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731

2 rows × 30 columns

In [21]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

In [22]:

```
x = cc_final.drop(['Class'], axis = 1)
y = cc_final['Class']
```

In [23]:

```
x.head()
```

Out[23]:

	Time	V1	V2	V3	V4	V5	V6	V7
270265	163994.0	-5.808461	5.392370	-5.774109	-0.357545	-2.558459	-1.157350	-3.267030
33592	37297.0	0.566593	-1.377684	0.438796	0.458684	-1.453851	-0.477246	-0.165218
107181	70320.0	-1.672614	-4.918784	-1.770582	1.282633	-1.757134	-0.065947	2.060316
146920	87962.0	-2.956152	2.569305	-1.167447	-2.839101	-0.504506	-1.310978	0.040740
154360	101272.0	1.263040	-1.097398	-0.603314	1.863709	-0.381287	0.407999	0.049589

5 rows × 30 columns

In [24]:

```
y.head()
```

Out[24]:

```
270265    0
33592     0
107181    0
146920    0
154360    0
Name: Class, dtype: int64
```

In [25]:

```
X_train, X_test, y_train, y_test = train_test_split (x, y, test_size=0.3, random_state=0)
```

In [26]:

```
lg = LogisticRegression()
```

In [27]:

```
lg.fit(X_train, y_train)
```

Out[27]:

```
LogisticRegression()
```

In [28]:

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[28]:

```
((688, 30), (296, 30), (688,), (296,))
```

In [29]:

```
y_train = y_train.values.reshape(-1,1)  
y_test = y_test.values.reshape(-1,1)
```

In [30]:

```
y_train.shape, y_test.shape
```

Out[30]:

```
((688, 1), (296, 1))
```

In [31]:

```
lg.score(X_test, y_test)
```

Out[31]:

```
0.9222972972972973
```


In [32]:

```
y_train_pred = lg.predict(X_train)
y_train_pred
```

Out[32]:

```
array([0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0,
       0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1,
       1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1,
       0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0,
       1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1,
       0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0,
       1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1,
       0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1,
       0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1,
       0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1,
       0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1,
       1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1,
       0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1,
       0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1,
       0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1,
       0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0,
       0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
       0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0,
       1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1,
       1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1,
       1, 1, 0, 1, 1, 1], dtype=int64)
```

In [33]:

```
y_test_pred = lg.predict(X_test)
y_test_pred
```

Out[33]:

```
array([0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1,
       1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0,
       1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1,
       0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0,
       0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0,
       1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0,
       0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1,
       1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0,
       1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1,
       0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0,
       1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1,
       0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0,
       1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0,
       0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1], dtype=int64)
```

In [34]:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

In [35]:

```
print("R2Score : ", r2_score(y_test, y_test_pred))
print("mean_absolute_error : ", mean_absolute_error(y_test, y_test_pred))
print("mean_squared_error : ", mean_squared_error(y_test, y_test_pred))
print("Root mean_squared_error : ", np.sqrt(mean_squared_error(y_test, y_test_pred)))
```

```
R2Score : 0.6891749988586038
mean_absolute_error : 0.0777027027027027
mean_squared_error : 0.0777027027027027
Root mean_squared_error : 0.27875204519913876
```

In [36]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train_std = sc.fit_transform(X_train)
X_test_std = sc.transform(X_test)
```

In [37]:

```
from sklearn.ensemble import RandomForestRegressor

rf_tree = RandomForestRegressor(random_state=0)
rf_tree.fit(X_train_std, y_train)
rf_tree_y_pred = rf_tree.predict(X_train_std)
print("Accuracy: {}".format(rf_tree.score(X_train_std, y_train)))
print("R squared: {}".format(r2_score(y_true=y_train, y_pred=rf_tree_y_pred)))
```

```
Accuracy: 0.9734927992563485
R squared: 0.9734927992563485
```

In [38]:

```
print ('Logistic Regression Accuracy: {}'.format(lg.score(X_test, y_test)))  
print ('Random Forest Accuracy: {}'.format(rf_tree.score(X_train_std,y_train)))
```

Logistic Regression Accuracy: 0.9222972972972973

Random Forest Accuracy: 0.9734927992563485

Conclusion

Random Forest Regressor has more accuracy than Logistic Regression. So we can use Random Forest Regressor to detect the fraud in credit transaction.